

Open in app ↗



Medium



Search



Write



Data Engineer Things



Member-only story

8 Awesome Docker Images Every Data Engineer Needs on Their Radar in 2025

By the end of this article, you'll be able to spin up production-ready environments for Airflow, Spark, Kafka, and more in seconds ,no more wrestling with OS packages or version conflicts.



Victor Oketch Sabare

5 min read · Jun 11, 2025




--



4





Why Docker Changed the Game for Data Engineers

Ever spent hours configuring Python, Java, and system libraries only to discover a subtle version mismatch in production? I certainly have. A few years ago, I lost half a day trying to reproduce an Airflow DAG on my colleague's Mac. Docker changed all that.

- **Portability:** Run the same image on your laptop, CI server, or Kubernetes cluster without tweaking dependencies.
- **Reproducibility:** Pin versions in your `Dockerfile` or `docker-compose.yml` so “it works on my machine” actually means **it works everywhere**.

- **Isolation:** Shield your host OS from conflicting library versions or experimental packages.


These benefits aren't just hype they're why I started running every component of my data platform in containers. And over time, I discovered a handful of community-approved Docker images that **truly** streamline data workflows. Here's my curated list of game-changers you need on your radar.

1. Apache Airflow — puckel/docker-airflow

I kicked off my container journey with this unofficial image. It bundles Airflow, Celery executors, and a Postgres backend — all in one.

```
docker pull puckel/docker-airflow:2.5.1
docker run -d -p 8080:8080 \
  -e LOAD_EX=n \
  -v ~/dags:/usr/local/airflow/dags \
  puckel/docker-airflow:2.5.1 webserver
```

- **Why it's game-changing:** Zero setup for the metadata database, worker pool, and web UI.
- **Pro tip:** Mount your DAG folder for live edits; add a `docker-compose.yml` to stand up the scheduler and worker in one command.




2. Apache Spark — bitnami/spark

For heavy lifting on large datasets, Spark is non-negotiable. The Bitnami image packs Hadoop and Spark in a ready-to-go cluster:

```
docker pull bitnami/spark:3  
docker run -it bitnami/spark spark-shell
```

- **My experience:** I spun up a 3-node mini-cluster on my laptop to prototype a real-time ETL pipeline in under 10 minutes.
- **Hot tip:** Use the companion `bitnami/hadoop` image for a full HDFS environment to test your code locally.




3. Apache Kafka — confluentinc/cp-kafka

Reliable streaming is crucial for event-driven pipelines. Confluent's CP-Kafka image comes battle-tested:

```
docker pull confluentinc/cp-kafka:7.4.0  
docker-compose up -d zookeeper kafka
```

- **Why it stood out:** Pre-configured defaults for topics, security, and Confluent Schema Registry.
- **Insider trick:** Wire up the Schema Registry image (`confluentinc/cp-schema-registry`) to enforce Avro schemas in your pipelines.




4. PostgreSQL — postgres:15

Although not exclusive to data engineering, Postgres remains the workhorse for metadata and analytical workloads:


```
docker pull postgres:15
docker run -d -p 5432:5432 \
  -e POSTGRES_PASSWORD=secret \
  -v db_data:/var/lib/postgresql/data \
  postgres:15
```

- **My shortcut:** Use `docker-entrypoint-initdb.d/` to bootstrap your analytics schema on first run.
- **Key benefit:** Built-in support for JSONB and extensions like PostGIS make Postgres a flexible choice for hybrid data stores.



5. JupyterLab — `jupyter/datascience-notebook`

Nothing beats the interactivity of Jupyter notebooks, and this official image bundles popular libraries:

```
docker pull jupyter/datascience-notebook:latest  
docker run -p 8888:8888 jupyter/datascience-notebook
```

- **Why I love it:** Immediate access to Pandas, Matplotlib, scikit-learn, PySpark, and more.
- **Best practice:** Attach your project folder as a volume so your notebooks and data persist between sessions.

6. MinIO — minio/minio

For local S3-compatible storage, MinIO is a lightweight, production-grade